

finalProjCode

May 3, 2025

1 Predicting Depression Risk Using NHANES Behavioral and Clinical Data: A Machine Learning Approach

1.1 *Matthew Maslow*

```
[3]: # Core Libraries
import os
import pandas as pd
import numpy as np

# File Reading
import pyreadstat

# Data Preprocessing
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.utils.class_weight import compute_class_weight
from imblearn.over_sampling import SMOTE

# Modeling
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier

# Evaluation
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay, roc_auc_score, roc_curve
from sklearn.model_selection import cross_val_score, validation_curve

# Visualization
import matplotlib.pyplot as plt
from xgboost import plot_importance

# Explainability
import shap

# Explainability
import shap
```

2 Load in Data

```
[5]: base_folder = "/Users/matthewmaslow/desktop/DS-596-Special Topics Medical_
    Science/Final Project"
data_folder = os.path.join(base_folder, "data")
output_folder = os.path.join(base_folder, "DataCSV")

selected_columns = {
    "mhDepressionScreener_nhanes.xpt": ["SEQN", "DPQ020", "DPQ060", "DPQ090"],
    "demographics_nhanes.xpt": ["SEQN", "RIAGENDR", "RIDAGEYR", "RIDRETH3", \
    DMDDEDUC2", "DMDMARTZ", "INDFMPIR"],
    "income_nhanes.xpt": ["SEQN", "INDFMMPI", "INDFMMPC", "INQ300", "IND310"],
```

```

    "sleepDisorders_nhanes.xpt": ["SEQN", "SLD012", "SLD013"],
    "smokingBehavior_nhanes.xpt": ["SEQN", "SMQ020", "SMQ040", "SMD650"],
    "alcoholUse_nhanes.xpt": ["SEQN", "ALQ111", "ALQ121", "ALQ270"],
    "physicalActivity_nhanes.xpt": ["SEQN", "PAD800", "PAD820", "PAD680"],
    "VID_L_nhanes.xpt": ["SEQN", "LBXVIDMS"],
    "GHB_L_nhanes.xpt": ["SEQN", "LBXGH"],
    "HSCRP_L_nhanes.xpt": ["SEQN", "LBXHSCRP"]
}

def load_selected_columns(file_path, columns):
    try:
        df, _ = pyreadstat.read_xport(file_path, usecols=columns)
        print(f"Loaded with pyreadstat: {os.path.basename(file_path)}")
    except Exception:
        df = pd.read_sas(file_path, format='xport', encoding='latin1')
        df = df[columns]
        print(f"Fallback to pandas: {os.path.basename(file_path)}")
    return df

merged_df = None
for file_name, cols in selected_columns.items():
    path = os.path.join(data_folder, file_name)
    df = load_selected_columns(path, cols)
    merged_df = df if merged_df is None else pd.merge(merged_df, df, on="SEQN",
↪how="left")

os.makedirs(output_folder, exist_ok=True)
output_path = os.path.join(output_folder, "nhanes_cleaned_merged.csv")
merged_df.to_csv(output_path, index=False)

print(f"\nMerged file saved to: {output_path}")
print(f"Final shape: {merged_df.shape}")

```

```

Loaded with pyreadstat: mhDepressionScreener_nhanes.xpt
Loaded with pyreadstat: demographics_nhanes.xpt
Loaded with pyreadstat: income_nhanes.xpt
Loaded with pyreadstat: sleepDisorders_nhanes.xpt
Loaded with pyreadstat: smokingBehavior_nhanes.xpt
Loaded with pyreadstat: alcoholUse_nhanes.xpt
Loaded with pyreadstat: physicalActivity_nhanes.xpt
Loaded with pyreadstat: VID_L_nhanes.xpt
Loaded with pyreadstat: GHB_L_nhanes.xpt
Loaded with pyreadstat: HSCRP_L_nhanes.xpt

```

```

Merged file saved to: /Users/matthewmaslow/desktop/DS-596-Special Topics Medical
Science/Final Project/DataCSV/nhanes_cleaned_merged.csv
Final shape: (6337, 28)

```

```
[6]: df = pd.read_csv(os.path.join(output_folder, "nhanes_cleaned_merged.csv"))

print(df.dtypes)
print(df.isnull().sum())
print(df.shape)
print(df.head())
```

```
SEQN          float64
DPQ020        float64
DPQ060        float64
DPQ090        float64
RIAGENDR      float64
RIDAGEYR      float64
RIDRETH3      float64
DMDEDUC2      float64
DMDMARTZ      float64
INDFMPIR      float64
INDFMMPI      float64
INDFMMPC      float64
INQ300        float64
IND310        float64
SLD012        float64
SLD013        float64
SMQ020        float64
SMQ040        float64
SMD650        float64
ALQ111        float64
ALQ121        float64
ALQ270        float64
PAD800        float64
PAD820        float64
PAD680        float64
LBXVIDMS      float64
LBXGH         float64
LBXHSCRP      float64
```

```
dtype: object
```

```
SEQN          0
DPQ020        819
DPQ060        827
DPQ090        831
RIAGENDR      0
RIDAGEYR      0
RIDRETH3      0
DMDEDUC2      273
DMDMARTZ      274
INDFMPIR      831
INDFMMPI      1286
INDFMMPC      520
```

```

INQ300      517
IND310     3476
SLD012       65
SLD013       67
SMQ020        2
SMQ040     3819
SMD650     5425
ALQ111      856
ALQ121     1415
ALQ270     3971
PAD800     1306
PAD820     3460
PAD680        6
LBXVIDMS    513
LBXGH       335
LBXHSCRIP   516

```

```
dtype: int64
```

```
(6337, 28)
```

	SEQN	DPQ020	DPQ060	DPQ090	RIAGENDR	RIDAGEYR	RIDRETH3	DMDEDUC2	\
0	130378.0	NaN	NaN	NaN	1.0	43.0	6.0	5.0	
1	130379.0	0.0	0.0	0.0	1.0	66.0	3.0	5.0	
2	130380.0	0.0	0.0	0.0	2.0	44.0	2.0	3.0	
3	130386.0	0.0	1.0	0.0	1.0	34.0	1.0	4.0	
4	130387.0	0.0	0.0	0.0	2.0	68.0	3.0	5.0	

	DMDMARTZ	INDFMPIR	...	SMD650	ALQ111	ALQ121	ALQ270	PAD800	PAD820	\
0	1.0	5.00	...	NaN	NaN	NaN	NaN	45.0	45.0	
1	1.0	5.00	...	NaN	1.0	2.0	NaN	45.0	45.0	
2	1.0	1.41	...	NaN	1.0	10.0	NaN	20.0	NaN	
3	1.0	1.33	...	NaN	1.0	4.0	0.0	30.0	30.0	
4	3.0	1.32	...	NaN	1.0	0.0	NaN	NaN	NaN	

	PAD680	LBXVIDMS	LBXGH	LBXHSCRIP
0	360.0	58.9	5.6	1.78
1	480.0	60.5	5.6	2.03
2	240.0	39.4	6.2	5.62
3	180.0	96.9	5.1	1.05
4	1200.0	26.7	5.9	3.96

```
[5 rows x 28 columns]
```

3 Data Preparation

```
[8]: df = df.dropna(subset=["DPQ020", "DPQ060", "DPQ090"])
```

```
[9]: df = df.drop(columns=["SMD650", "SMQ040", "ALQ270", "PAD820", "IND310"])
```

```
[10]: num_cols = df.select_dtypes(include="number").columns.tolist()
num_cols = [col for col in num_cols if col not in ["SEQN", "DPQ020", "DPQ060",
↪ "DPQ090"]]

imputer = SimpleImputer(strategy="median")
df[num_cols] = imputer.fit_transform(df[num_cols])

[11]: invalid_vals = [7, 9, 77, 99, 777, 7777, 999, 9999]

df = df[~df.isin(invalid_vals).any(axis=1)]

df.reset_index(drop=True, inplace=True)

[12]: df['DPQ020_binary'] = df['DPQ020'].apply(lambda x: 0 if x == 0 else 1)
df['DPQ060_binary'] = df['DPQ060'].apply(lambda x: 0 if x == 0 else 1)
```

4 Model Training

4.0.1 DPQ020 – Feeling down, depressed, or hopeless

```
[15]: classes = np.array(sorted(df["DPQ020"].unique()))
class_weights_array = compute_class_weight(class_weight='balanced',
↪ classes=classes, y=df["DPQ020"])
class_weights = dict(zip(classes, class_weights_array))
print(class_weights_array)

[0.37325508 1.07149362 4.74395161 5.65625   ]

[16]: X = df.drop(columns=["SEQN", "DPQ020", "DPQ060", "DPQ090", "DPQ020_binary",
↪ "DPQ060_binary"])
y = df["DPQ020_binary"]

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
↪ test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

classes = np.unique(y_train)
weights = compute_class_weight(class_weight='balanced', classes=classes,
↪ y=y_train)
class_weights = dict(zip(classes, weights))

log_reg = LogisticRegression(solver='lbfgs', max_iter=1000,
↪ class_weight=class_weights)
log_reg.fit(X_train_scaled, y_train)
```

```

y_pred_log = log_reg.predict(X_test_scaled)

print("Logistic Regression Results (DPQ020_binary)")
print(classification_report(y_test, y_pred_log))
print(confusion_matrix(y_test, y_pred_log))

explainer_log = shap.Explainer(log_reg, X_train_scaled, feature_names=X.columns)
shap_values_log = explainer_log(X_test_scaled)
plt.figure()
shap.summary_plot(shap_values_log, X_test, plot_type="bar", show=False)
plt.title("SHAP Feature Importance (Logistic Regression - DPQ020_binary)")
plt.tight_layout()
plt.show()

xgb = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    learning_rate=0.05,
    n_estimators=1000,
    max_depth=10,
    scale_pos_weight=class_weights[1] / class_weights[0]
)

xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)

print("XGBoost Results (DPQ020_binary)")
print(classification_report(y_test, y_pred_xgb))
print(confusion_matrix(y_test, y_pred_xgb))

explainer_xgb = shap.Explainer(xgb)
shap_values_xgb = explainer_xgb(X_test)
plt.figure()
shap.summary_plot(shap_values_xgb, X_test, plot_type="bar", show=False)
plt.title("SHAP Feature Importance (XGBoost - DPQ020_binary)")
plt.tight_layout()
plt.show()

```


```

Logistic Regression Results (DPQ020_binary)

```

	precision	recall	f1-score	support
0	0.79	0.59	0.68	315
1	0.45	0.69	0.55	156
accuracy			0.62	471
macro avg	0.62	0.64	0.61	471
weighted avg	0.68	0.62	0.63	471

```
[[186 129]
 [ 49 107]]
```




XGBoost Results (DPQ020_binary)

	precision	recall	f1-score	support
0	0.72	0.80	0.76	315
1	0.49	0.38	0.43	156

accuracy			0.66	471
macro avg	0.61	0.59	0.60	471
weighted avg	0.65	0.66	0.65	471

```
[[252 63]
 [ 96 60]]
```



Validation

```
[18]: log_scores = cross_val_score(log_reg, X_train_scaled, y_train, cv=5,
    ↪scoring='f1')
print("Logistic Regression Cross-Validated F1 Score (mean ± std):", log_scores.
    ↪mean(), "+/-", log_scores.std())

xgb_scores = cross_val_score(xgb, X_train, y_train, cv=5, scoring='f1')
print("XGBoost Cross-Validated F1 Score (mean ± std):", xgb_scores.mean(), "+/"
    ↪"-", xgb_scores.std())
```

Logistic Regression Cross-Validated F1 Score (mean ± std): 0.5348596434072617
 +/- 0.02329572659756699
 XGBoost Cross-Validated F1 Score (mean ± std): 0.3990104656138901 +/-
 0.018777327363326997

```
[19]: param_range_logreg = np.logspace(-4, 2, 6)
train_scores_log, test_scores_log = validation_curve(
    LogisticRegression(max_iter=1000),
    X_train_scaled,
    y_train,
    param_name='C',
    param_range=param_range_logreg,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)

param_range_xgb = [3, 5, 7, 9, 11]
train_scores_xgb, test_scores_xgb = validation_curve(
    XGBClassifier(learning_rate=0.05, n_estimators=1000, eval_metric='logloss'),
    X_train,
    y_train,
    param_name='max_depth',
    param_range=param_range_xgb,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)

fig, axs = plt.subplots(1, 2, figsize=(14, 5))

axs[0].plot(param_range_logreg, train_scores_log.mean(axis=1), label="Train",
    ↪marker='o')
axs[0].plot(param_range_logreg, test_scores_log.mean(axis=1),
    ↪label="Validation", marker='o')
axs[0].set_xscale('log')
axs[0].set_xlabel("C (Inverse Regularization)")
axs[0].set_ylabel("Accuracy")
```


```

axs[0].set_title("Validation Curve - Logistic Regression")
axs[0].legend()
axs[0].grid(True)

axs[1].plot(param_range_xgb, train_scores_xgb.mean(axis=1), label="Train",
            ↪marker='o')
axs[1].plot(param_range_xgb, test_scores_xgb.mean(axis=1), label="Validation",
            ↪marker='o')
axs[1].set_xlabel("Max Depth")
axs[1].set_ylabel("Accuracy")
axs[1].set_title("Validation Curve - XGBoost")
axs[1].legend()
axs[1].grid(True)

plt.tight_layout()
plt.show()

```



4.0.2 DPQ060_binary - Feeling bad about yourself

```

[21]: classes = np.array(sorted(df["DPQ060"].unique()))
      class_weights_array = compute_class_weight(class_weight='balanced',
      ↪classes=classes, y=df["DPQ060"])
      class_weights = dict(zip(classes, class_weights_array))
      print(class_weights_array)

[0.34200581  1.34919725  5.29954955  6.84011628]

[22]: X = df.drop(columns=["SEQN", "DPQ020", "DPQ060", "DPQ090", "DPQ020_binary",
      ↪"DPQ060_binary"])
      y = df["DPQ060_binary"]

      X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
      ↪test_size=0.2, random_state=42)

```

```

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

classes = np.unique(y_train)
weights = compute_class_weight(class_weight='balanced', classes=classes,
    ↪y=y_train)
class_weights = dict(zip(classes, weights))

log_reg = LogisticRegression(solver='lbfgs', max_iter=1000,
    ↪class_weight=class_weights)
log_reg.fit(X_train_scaled, y_train)
y_pred_log = log_reg.predict(X_test_scaled)

print("Logistic Regression Results (DPQ060_binary)")
print(classification_report(y_test, y_pred_log))
print(confusion_matrix(y_test, y_pred_log))

explainer_log = shap.Explainer(log_reg, X_train_scaled, feature_names=X.columns)
shap_values_log = explainer_log(X_test_scaled)
plt.figure()
shap.summary_plot(shap_values_log, X_test, plot_type="bar", show=False)
plt.title("SHAP Feature Importance (Logistic Regression - DPQ060_binary)")
plt.tight_layout()
plt.show()

xgb = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    learning_rate=0.05,
    n_estimators=1000,
    max_depth=10,
    scale_pos_weight=class_weights[1] / class_weights[0]
)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)

print("XGBoost Results (DPQ060_binary)")
print(classification_report(y_test, y_pred_xgb))
print(confusion_matrix(y_test, y_pred_xgb))

explainer_xgb = shap.Explainer(xgb)
shap_values_xgb = explainer_xgb(X_test)
plt.figure()
shap.summary_plot(shap_values_xgb, X_test, plot_type="bar", show=False)
plt.title("SHAP Feature Importance (XGBoost - DPQ060_binary)")


```

```
plt.tight_layout()
plt.show()
```

Logistic Regression Results (DPQ060_binary)

	precision	recall	f1-score	support
0	0.84	0.66	0.74	344
1	0.41	0.65	0.50	127
accuracy			0.66	471
macro avg	0.62	0.65	0.62	471
weighted avg	0.72	0.66	0.68	471


```
[[228 116]
 [ 45  82]]
```



XGBoost Results (DPQ060_binary)

	precision	recall	f1-score	support
0	0.76	0.84	0.80	344
1	0.39	0.27	0.32	127
accuracy			0.69	471
macro avg	0.57	0.56	0.56	471
weighted avg	0.66	0.69	0.67	471

```
[[290 54]
 [ 93 34]]
```



Validation

```
[24]: log_scores = cross_val_score(log_reg, X_train_scaled, y_train, cv=5,
    ↪scoring='f1')
```

```

print("Logistic Regression Cross-Validated F1 Score (mean ± std):", log_scores.
    ↪mean(), "+/-", log_scores.std())

xgb_scores = cross_val_score(xgb, X_train, y_train, cv=5, scoring='f1')
print("XGBoost Cross-Validated F1 Score (mean ± std):", xgb_scores.mean(), "+/"
    ↪ "-", xgb_scores.std())

```

Logistic Regression Cross-Validated F1 Score (mean ± std): 0.4679322508665883
 +/- 0.018278261622465264
 XGBoost Cross-Validated F1 Score (mean ± std): 0.34185807295920434 +/-
 0.010802380228495638

```

[25]: param_range_logreg = np.logspace(-4, 2, 6)
train_scores_log, test_scores_log = validation_curve(
    LogisticRegression(max_iter=1000),
    X_train_scaled,
    y_train,
    param_name='C',
    param_range=param_range_logreg,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)

param_range_xgb = [3, 5, 7, 9, 11]
xgb = XGBClassifier(
    learning_rate=0.05,
    n_estimators= 1000,
    early_stopping_rounds=10,
    max_depth=10,
    use_label_encoder=False,
    eval_metric='logloss'
)

fig, axs = plt.subplots(1, 2, figsize=(14, 5))

axs[0].plot(param_range_logreg, train_scores_log.mean(axis=1), label="Train",
    ↪marker='o')
axs[0].plot(param_range_logreg, test_scores_log.mean(axis=1),
    ↪label="Validation", marker='o')
axs[0].set_xscale('log')
axs[0].set_xlabel("C (Inverse Regularization)")
axs[0].set_ylabel("Accuracy")
axs[0].set_title("Validation Curve - Logistic Regression")
axs[0].legend()
axs[0].grid(True)

```




```

axs[1].plot(param_range_xgb, train_scores_xgb.mean(axis=1), label="Train",
            ↪marker='o')
axs[1].plot(param_range_xgb, test_scores_xgb.mean(axis=1), label="Validation",
            ↪marker='o')
axs[1].set_xlabel("Max Depth")
axs[1].set_ylabel("Accuracy")
axs[1].set_title("Validation Curve - XGBoost")
axs[1].legend()
axs[1].grid(True)

plt.tight_layout()
plt.show()

```



4.0.3 Imbalanced Learning (SMOTE)

```

[27]: def train_smote_models(X, y, label):
    print(f"\n=====")
    print(f"Running Models for {label}")
    print(f"=====")

    X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
    ↪test_size=0.2, random_state=42)

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    sm = SMOTE(random_state=42)
    X_train_res, y_train_res = sm.fit_resample(X_train_scaled, y_train)

    log_reg = LogisticRegression(solver='lbfgs', max_iter=1000, C=1.0,
    ↪class_weight=None)

```

```

log_reg.fit(X_train_res, y_train_res)
y_pred_log = log_reg.predict(X_test_scaled)

print(f"\nLogistic Regression Results - {label}")
print(classification_report(y_test, y_pred_log))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_log))

xgb = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    learning_rate=0.05,
    n_estimators=1000,
    max_depth=10
)
xgb.fit(X_train_res, y_train_res)
y_pred_xgb = xgb.predict(X_test_scaled)

print(f"\nXGBoost Results - {label}")
print(classification_report(y_test, y_pred_xgb))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_xgb))

explainer_xgb = shap.Explainer(xgb)
shap_values_xgb = explainer_xgb(X_test)
plt.figure()
shap.summary_plot(shap_values_xgb, X_test, plot_type="bar", show=False)
plt.title(f"SHAP Feature Importance (XGBoost - {label})")
plt.tight_layout()
plt.show()

# Validation Curves
param_range_logreg = np.logspace(-4, 2, 6)
train_scores_log, test_scores_log = validation_curve(
    LogisticRegression(max_iter=1000),
    X_train_res,
    y_train_res,
    param_name='C',
    param_range=param_range_logreg,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)

param_range_xgb = [3, 5, 7, 9, 11]
train_scores_xgb, test_scores_xgb = validation_curve(
    XGBClassifier(learning_rate=0.05, n_estimators=1000,
eval_metric='logloss'),

```

```

        X_train_res,
        y_train_res,
        param_name='max_depth',
        param_range=param_range_xgb,
        cv=5,
        scoring='accuracy',
        n_jobs=-1
    )

    fig, axs = plt.subplots(1, 2, figsize=(14, 5))

    axs[0].plot(param_range_logreg, train_scores_log.mean(axis=1),
        label="Train", marker='o')
    axs[0].plot(param_range_logreg, test_scores_log.mean(axis=1),
        label="Validation", marker='o')
    axs[0].set_xscale('log')
    axs[0].set_xlabel("C (Inverse Regularization)")
    axs[0].set_ylabel("Accuracy")
    axs[0].set_title(f"Validation Curve - Logistic ({label})")
    axs[0].legend()
    axs[0].grid(True)

    axs[1].plot(param_range_xgb, train_scores_xgb.mean(axis=1), label="Train",
        marker='o')
    axs[1].plot(param_range_xgb, test_scores_xgb.mean(axis=1),
        label="Validation", marker='o')
    axs[1].set_xlabel("Max Depth")
    axs[1].set_ylabel("Accuracy")
    axs[1].set_title(f"Validation Curve - XGBoost ({label})")
    axs[1].legend()
    axs[1].grid(True)

    plt.tight_layout()
    plt.show()

```

```

[28]: X_020 = df.drop(columns=["SEQN", "DPQ020", "DPQ060", "DPQ090", "DPQ020_binary",
        ↪ "DPQ060_binary"])
y_020 = df["DPQ020_binary"]
train_smote_models(X_020, y_020, "DPQ020_binary")

X_060 = df.drop(columns=["SEQN", "DPQ020", "DPQ060", "DPQ090", "DPQ020_binary",
        ↪ "DPQ060_binary"])
y_060 = df["DPQ060_binary"]
train_smote_models(X_060, y_060, "DPQ060_binary")

```

=====

Running Models for DPQ020_binary

=====

Logistic Regression Results - DPQ020_binary

	precision	recall	f1-score	support
0	0.78	0.58	0.67	315
1	0.44	0.66	0.53	156
accuracy			0.61	471
macro avg	0.61	0.62	0.60	471
weighted avg	0.67	0.61	0.62	471

Confusion Matrix:

[[184 131]

[53 103]]


XGBoost Results - DPQ020_binary

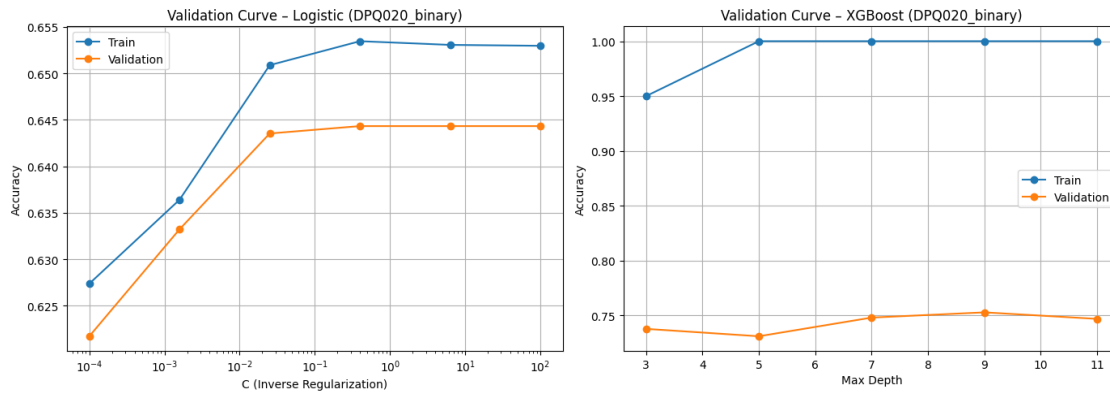
	precision	recall	f1-score	support
0	0.70	0.77	0.73	315
1	0.41	0.32	0.36	156
accuracy			0.62	471
macro avg	0.55	0.54	0.54	471
weighted avg	0.60	0.62	0.61	471

Confusion Matrix:

[[242 73]

[106 50]]





```
=====
Running Models for DPQ060_binary
=====
```

Logistic Regression Results - DPQ060_binary

	precision	recall	f1-score	support
0	0.84	0.65	0.73	344
1	0.41	0.65	0.50	127
accuracy			0.65	471
macro avg	0.62	0.65	0.62	471
weighted avg	0.72	0.65	0.67	471

Confusion Matrix:

```
[[223 121]
 [ 44  83]]
```

XGBoost Results - DPQ060_binary

	precision	recall	f1-score	support
0	0.76	0.83	0.79	344
1	0.38	0.28	0.33	127
accuracy			0.68	471
macro avg	0.57	0.56	0.56	471
weighted avg	0.66	0.68	0.67	471

Confusion Matrix:

```
[[286  58]
 [ 91  36]]
```

